

Regressione lineare con regolarizzazione L_2

Consideriamo il modello di **regressione lineare**, con un termine di regolarizzazione L_2 (**Ridge regression**)

$$\begin{aligned} E_{\text{aug}}(\boldsymbol{\theta}) \equiv J(\boldsymbol{\theta}) &= \frac{1}{N} \sum_{i=1}^N (y(i) - \boldsymbol{\varphi}^\top(i) \boldsymbol{\theta})^2 + \lambda_{\text{reg}} \cdot \sum_{j=0}^{d-1} (\theta_j)^2 \\ &= \frac{1}{N} \|Y - X \cdot \boldsymbol{\theta}\|_2^2 + \lambda_{\text{reg}} \cdot \|\boldsymbol{\theta}\|_2^2 \end{aligned}$$

Si può dimostrare che la stima in forma chiusa dei parametri si ottiene come:

$$\widehat{\boldsymbol{\theta}}_{\text{reg}} = \left(X^\top X + \lambda_{\text{reg}} \cdot I_d \right)^{-1} X^\top Y$$



Ridge regression e gradient descent

Supponiamo di avere **solo 2 parametri** $\theta = [\theta_0, \theta_1]^\top \in \mathbb{R}^{2 \times 1}$ per semplicità

$$E_{\text{aug}}(\theta) \equiv J(\theta) = \frac{1}{N} \sum_{i=1}^N (y(i) - \theta_0 - \theta_1 \cdot \varphi_1(i))^2 + \lambda_{\text{reg}} \cdot \sum_{j=0}^{d-1} (\theta_j)^2$$

E' possibile implementare l'algoritmo del gradient descent come:

For {

$$\theta_0 = \theta_0 - \alpha \cdot 2 \left[\frac{1}{N} \sum_{i=1}^N (y(i) - \theta_0 - \theta_1 \cdot \varphi_1(i)) \cdot (-1) + \lambda_{\text{reg}} \cdot \theta_0 \right]$$

$$\theta_1 = \theta_1 - \alpha \cdot 2 \left[\frac{1}{N} \sum_{i=1}^N (y(i) - \theta_0 - \theta_1 \cdot \varphi_1(i)) \cdot (-\varphi_1(i)) + \lambda_{\text{reg}} \cdot \theta_1 \right]$$

}



Regressione logistica con regolarizzazione L_2

Consideriamo il modello di **regressione logistica**, con un termine di regolarizzazione L_2

$$E_{\text{aug}}(\boldsymbol{\theta}) \equiv J(\boldsymbol{\theta}) = \sum_{i=1}^N \left(y(i) \cdot \ln \pi(i; \boldsymbol{\theta}) + (1 - y(i)) \cdot \ln[1 - \pi(i; \boldsymbol{\theta})] \right) + \lambda_{\text{reg}} \cdot \sum_{j=0}^{d-1} (\theta_j)^2$$

E' possibile implementare l'algoritmo del gradient descent come:

For {

$$\theta_0 = \theta_0 - \alpha \cdot \sum_{i=1}^N 1 \cdot (\pi(i) - y(i)) + 2\lambda_{\text{reg}} \cdot \theta_0$$

⋮

$$\theta_{d-1} = \theta_{d-1} - \alpha \cdot \sum_{i=1}^N \varphi_{d-1}(i) \cdot (\pi(i) - y(i)) + 2\lambda_{\text{reg}} \cdot \theta_{d-1}$$

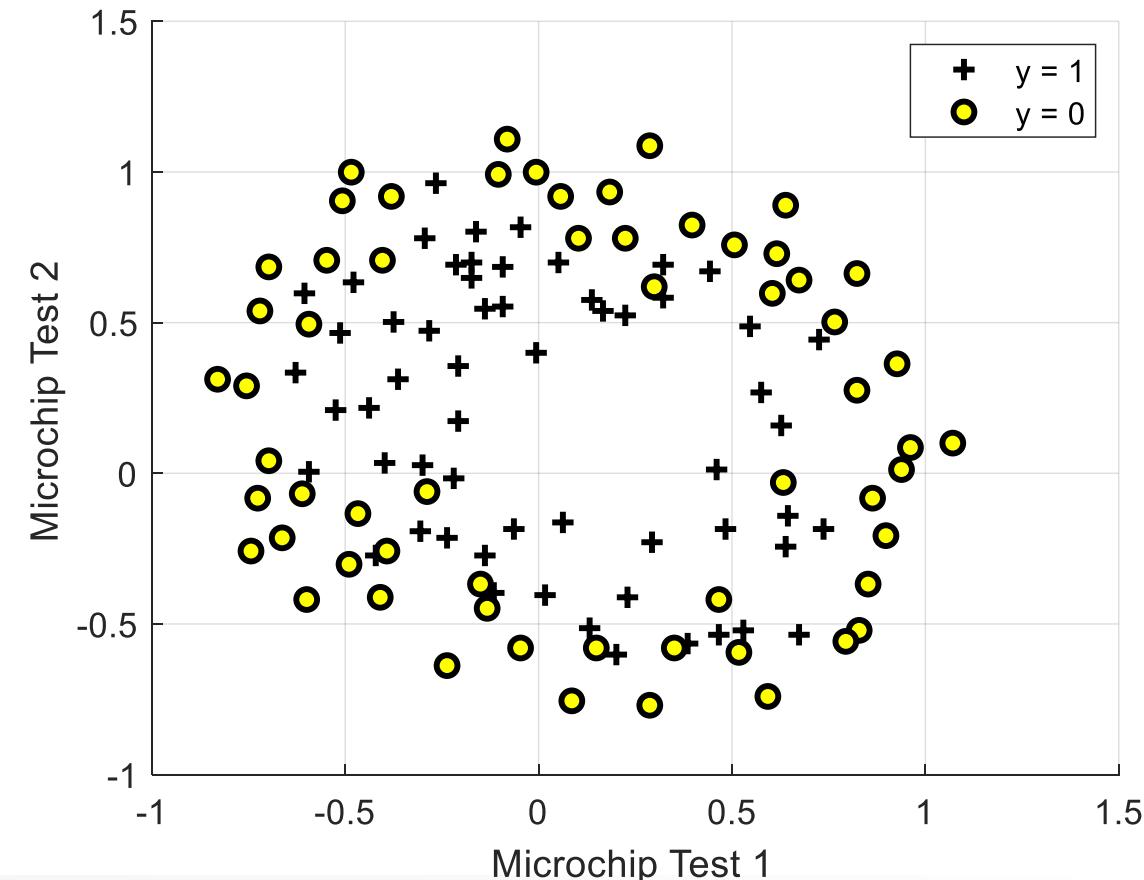
$$\begin{aligned} \pi(i) &\equiv \frac{1}{1 + e^{-\boldsymbol{\varphi}^\top(i)\boldsymbol{\theta}}} \\ &= P(y(i) = 1 | \boldsymbol{\varphi}(i)) \end{aligned}$$



Esercizio: classificare microchips difettosi

Vogliamo rilevare se un microchip è **difettoso** in base ai risultati di due test di qualità alla fine della linea di produzione, tramite un modello di **regressione logistica**

- Ogni microchip è descritto dalle seguenti features
 - ✓ φ_1 : Risultato del test 1
 - ✓ φ_2 : Risultato del test 2
- Il dataset è costituito da $N = 118$ microchips

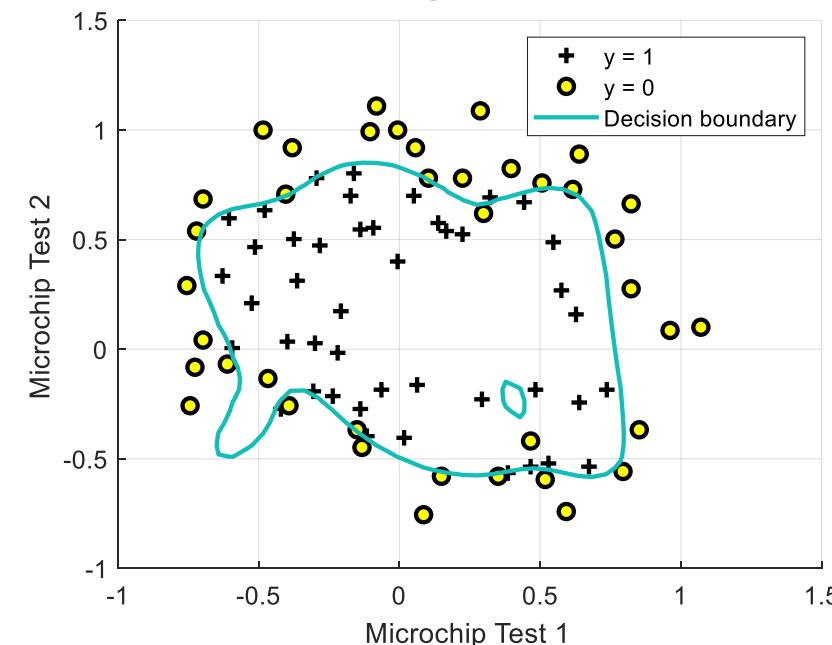


Esercizio: classificare microchips difettosi

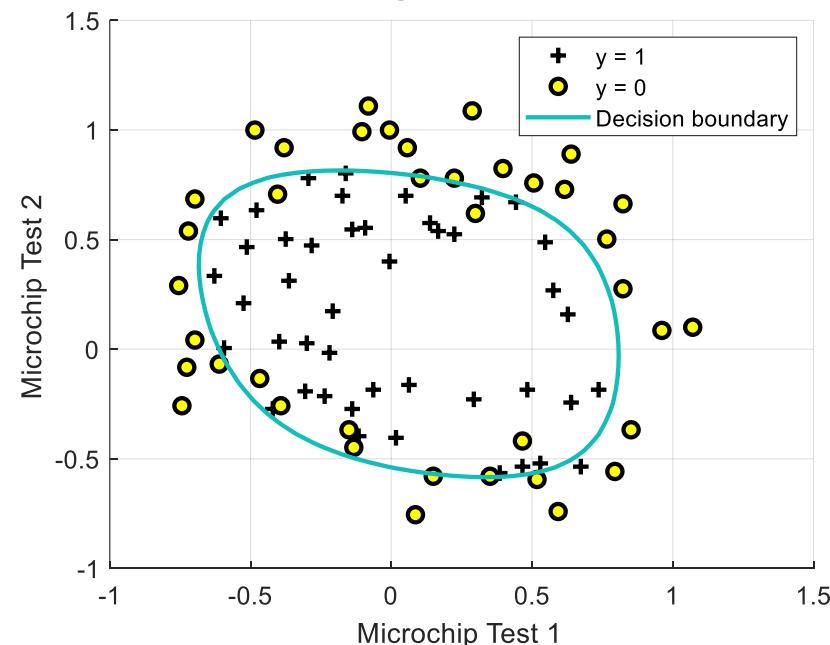
Per ottenere un **confine nonlineare** tramite il classificatore lineare, è possibile usare **features polinomiali**. Ad esempio, usando feature polinomiali di grado 2 otteniamo:

$$\varphi_1^2, \varphi_2^2, \dots, \varphi_{d-1}^2, \quad \varphi_1\varphi_2, \dots, \varphi_1\varphi_{d-1}, \quad \varphi_1\varphi_2^2, \dots, \varphi_1\varphi_{d-1}^2, \quad \varphi_1^2\varphi_2, \dots, \varphi_1^2\varphi_{d-1}, \dots$$

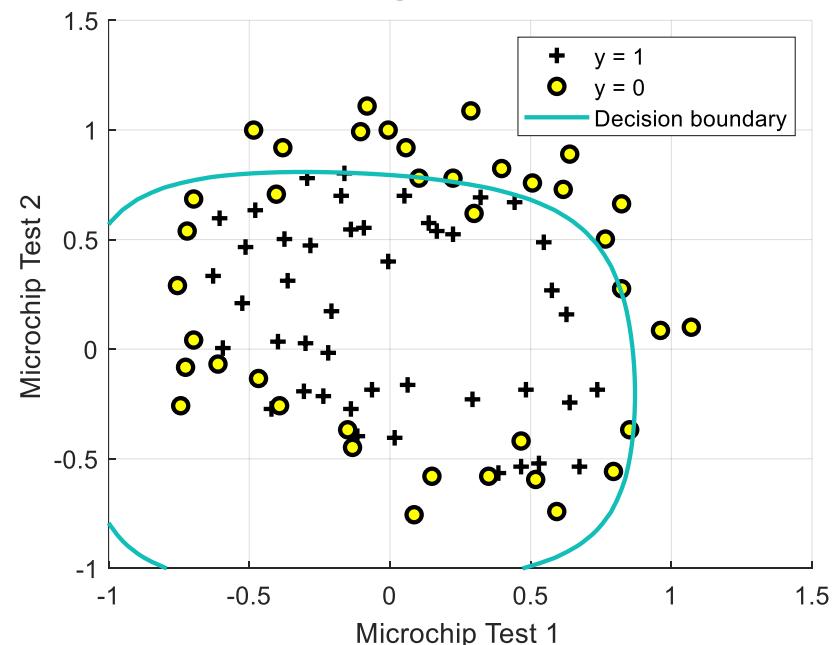
$$\lambda_{\text{reg}} = 0$$



$$\lambda_{\text{reg}} = 0.1$$

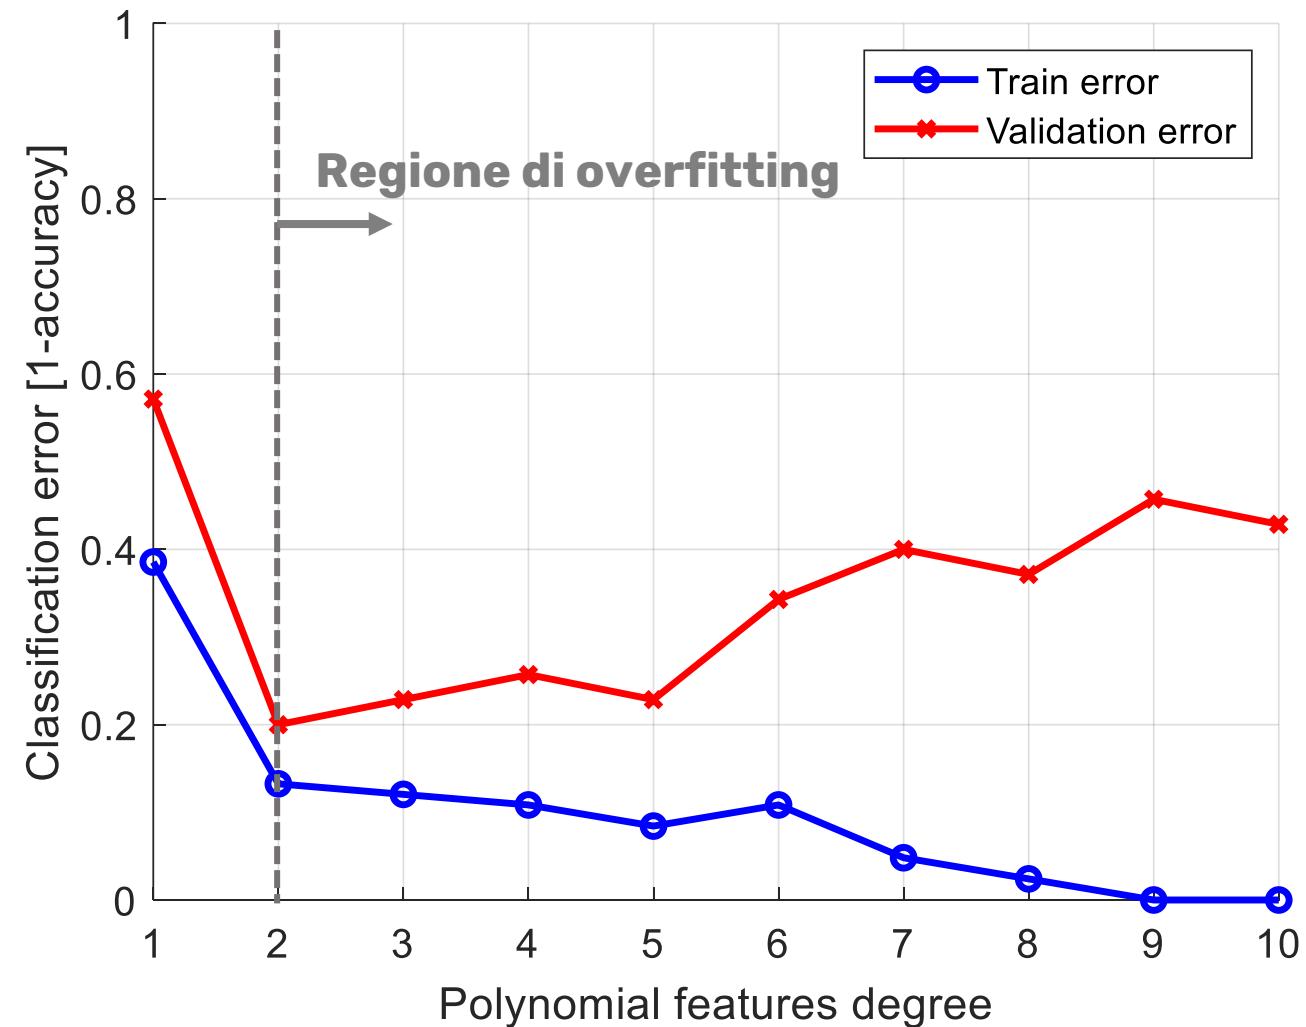


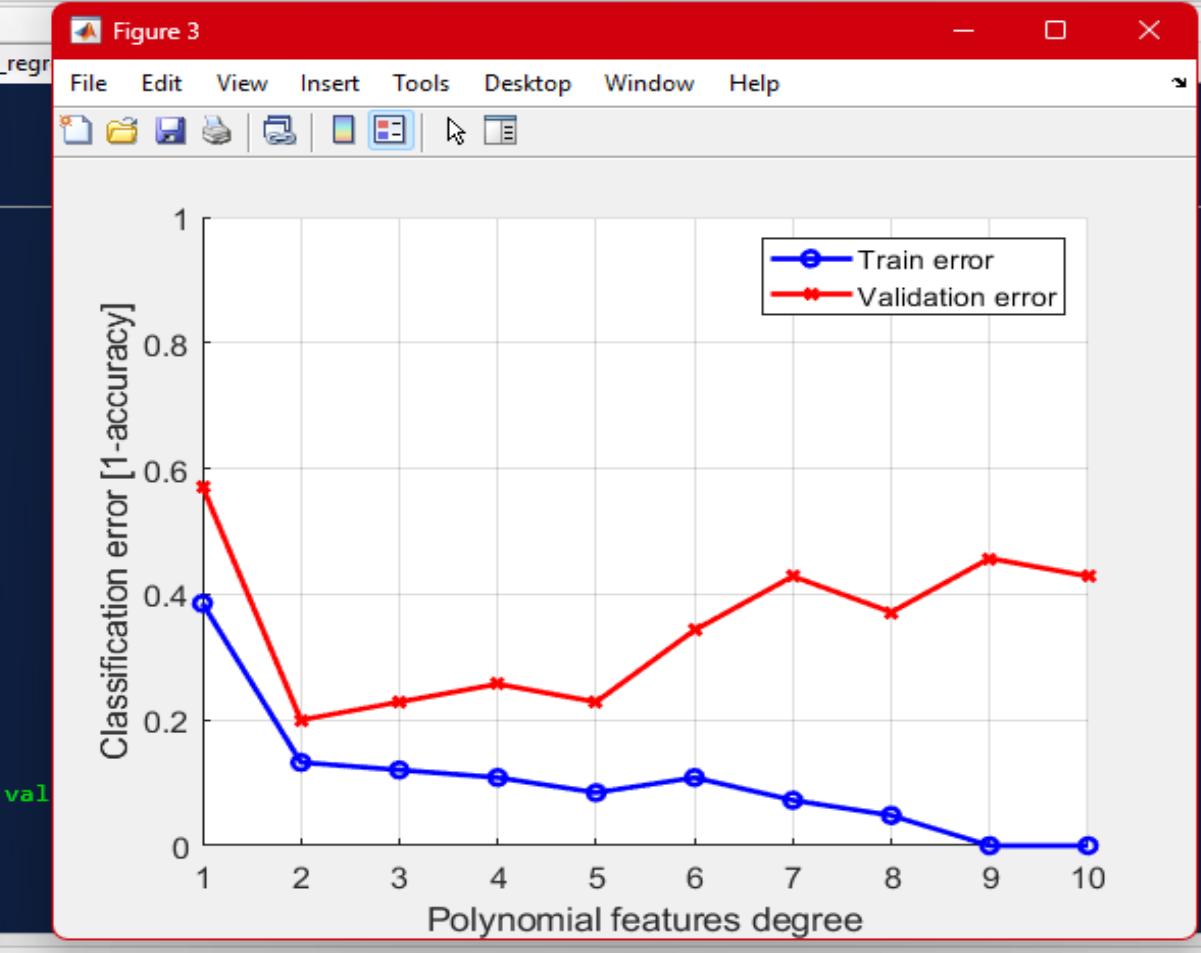
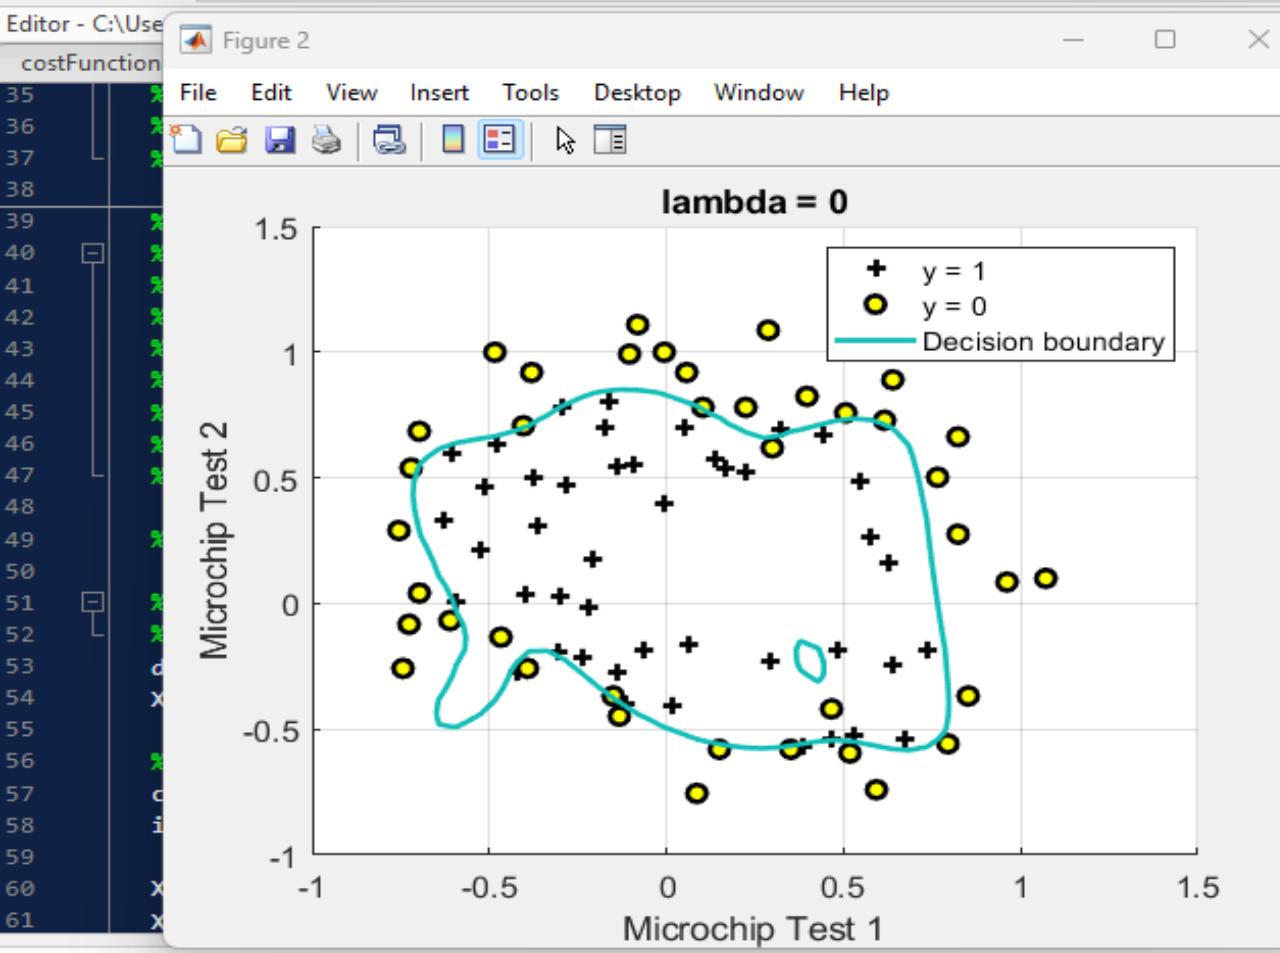
$$\lambda_{\text{reg}} = 10$$



Esercizio: classificare microchips difettosi

È possibile dividere i dati in training e validation set, in modo da **selezionare l'ordine ottimale** per le features polinomiali tramite **validazione**





Cost at initial theta (zeros): 57.531216

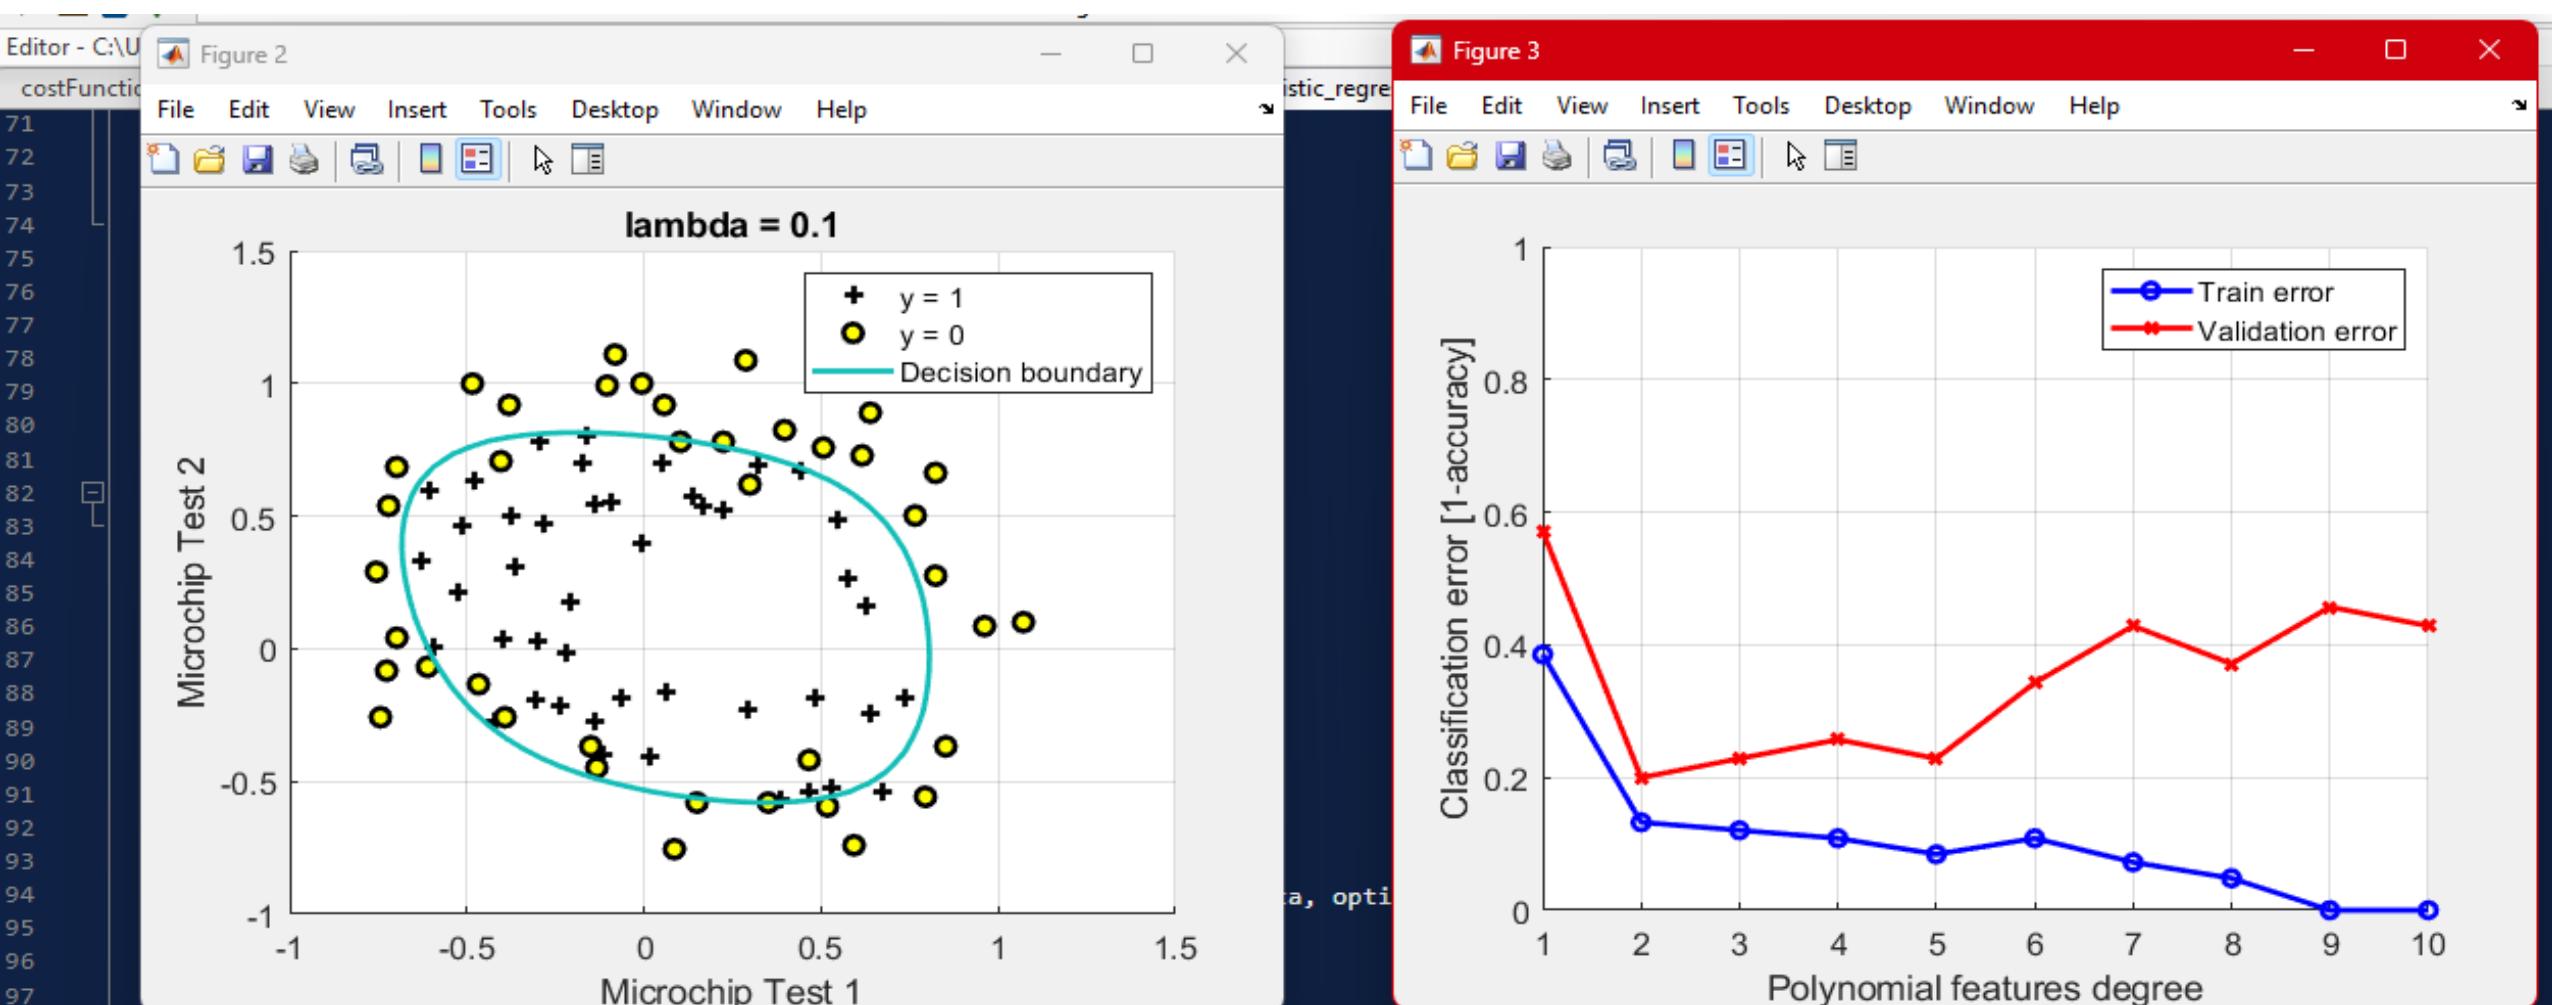
Local minimum possible.

fminunc stopped because it cannot decrease the objective function along the current search direction.

<stopping criteria details>

Train Accuracy: 91.566265

Validation Accuracy: 71.428571



Command Window

Cost at initial theta (zeros): 57.531216

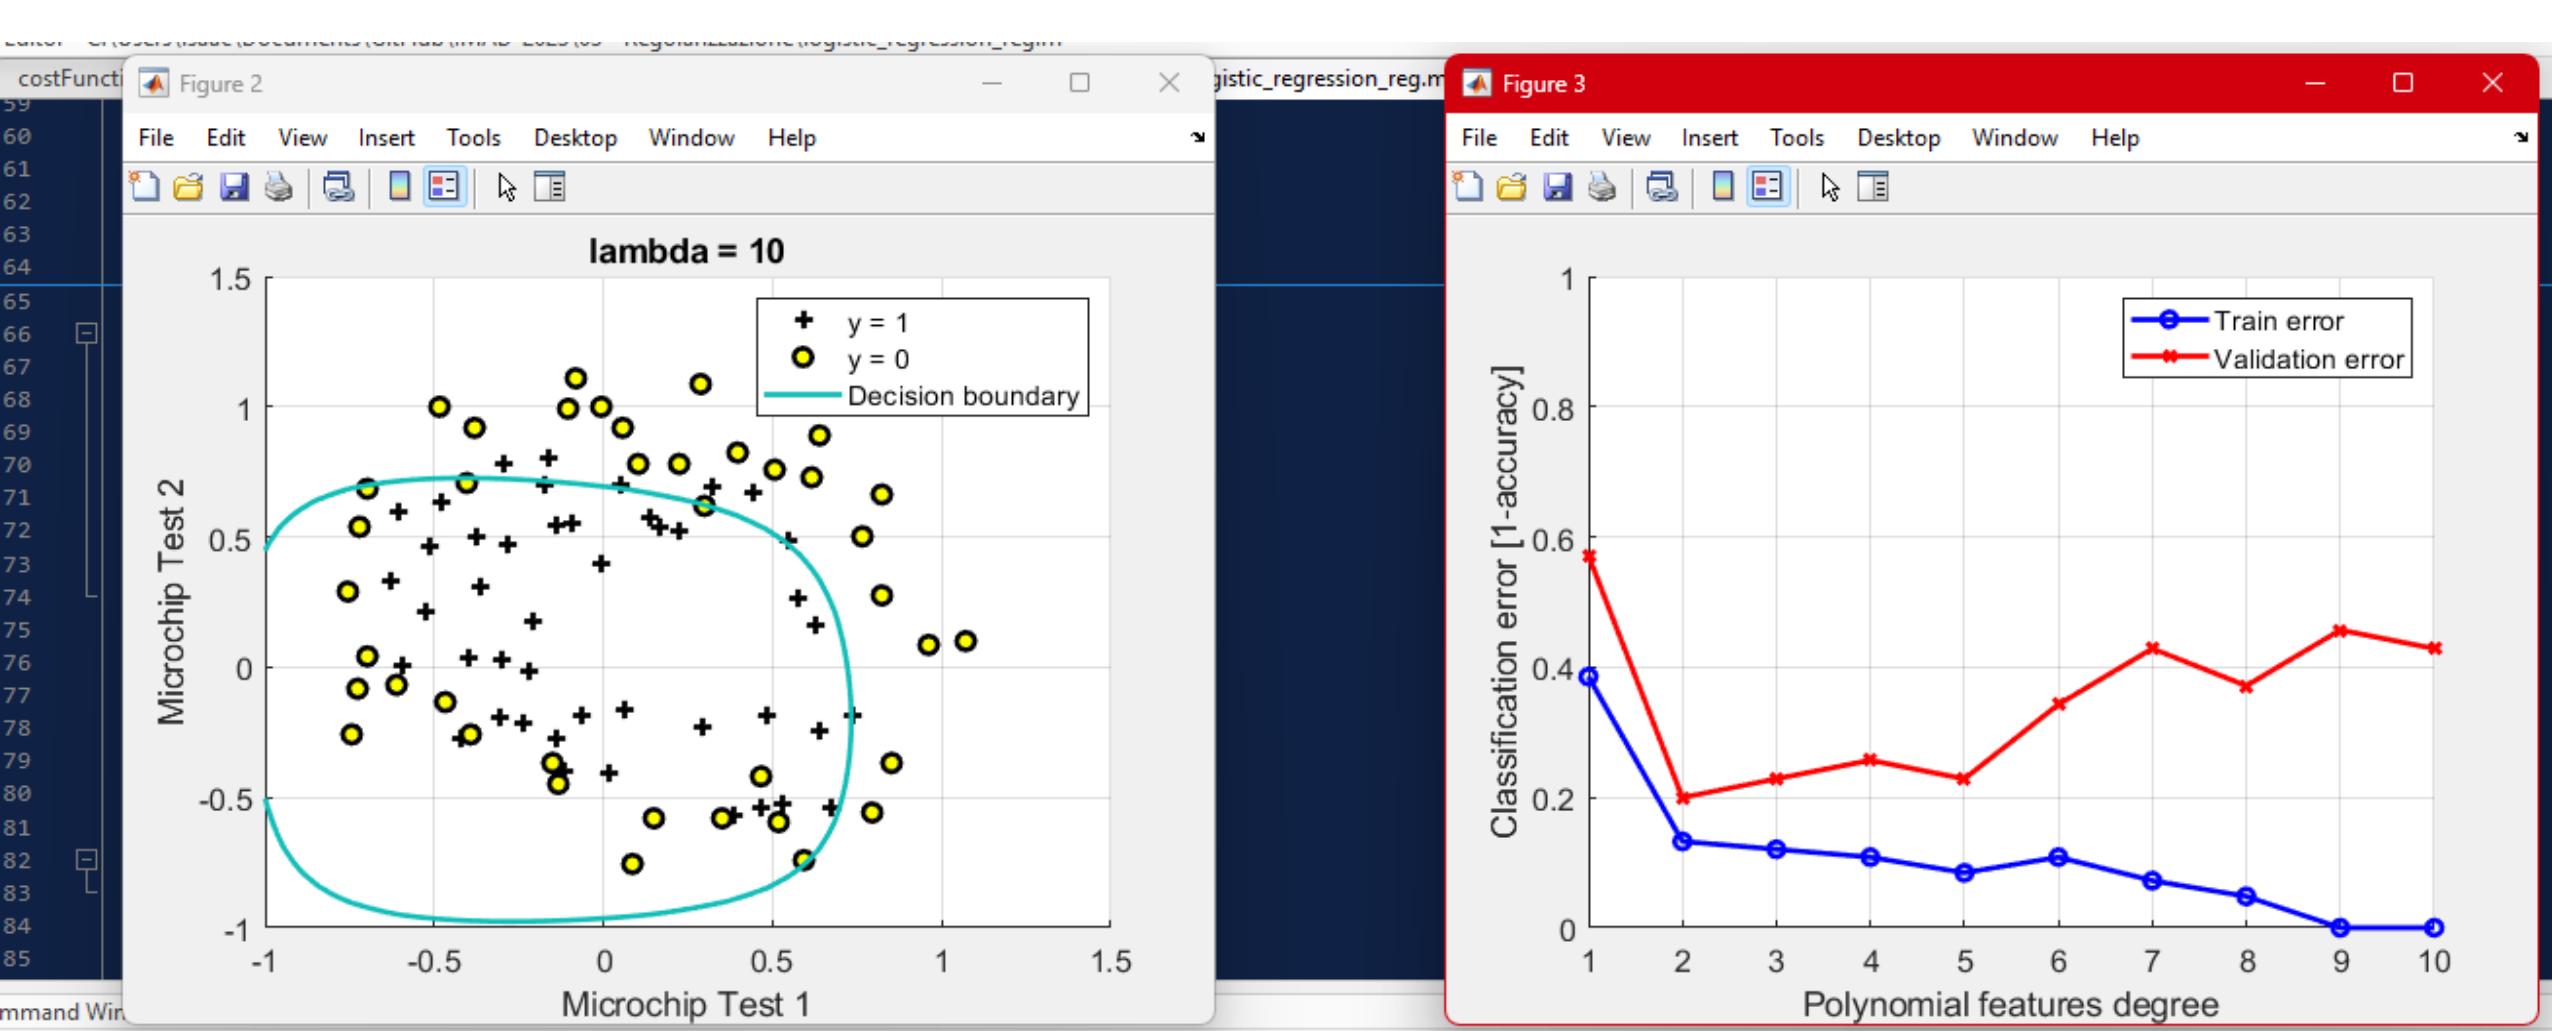
Local minimum possible.

fminunc stopped because it cannot decrease the objective function along the current search direction.

<stopping criteria details>

Train Accuracy: 87.951807

Validation Accuracy: 77.142857



Cost at initial theta (zeros): 57.531216

Local minimum possible.

fminunc stopped because it cannot decrease the objective function along the current search direction.

<stopping criteria details>

Train Accuracy: 68.674699

Validation Accuracy: 57.142857